

DAY 1: RSTAN, STATISTICAL MODELS

---

# STAN WORKSHOP

# POISSON DATA

# POISSON.DATA

---

► R:  
`source("poisson.data.R")`

► Looks like:

```
N <- 1000
```

```
x <- c(4,7,2,1,0,1,5,0,2,0,5,6,7,9, ...
```

# STEP 1: LOOK AT DATA; BUILD POISSON MODEL

---

- ▶ Hint:

```
data {  
  int N;  
  int<lower=0> x[N];  
}
```

$$x_n \sim \text{poisson}(\lambda)$$

or equivalently

- ▶ Also write a posterior predictive check:

```
generated quantities {  
  int x_ppc;  
  x_ppc <- poisson_rng(lambda);  
}
```

$$\begin{aligned} p(\lambda, x) &= \prod_{n=1}^N \frac{1}{x_n!} \lambda^{x_n} \exp(-\lambda) \\ &= \exp \left( \sum_{n=1}^N \text{poisson\_log}(x_n, \lambda) \right) \end{aligned}$$

- ▶ Try to do this using `increment_log_prob()` directly (you'll see why later)

---

## STEP 1: LOOK AT DATA; BUILD POISSON MODEL

```
data {  
  int N;  
  real x[N];  
}  
parameters {  
  real<lower = 0> lambda;  
}  
model {  
  lambda ~ cauchy(0, 1);  
  x ~ poisson(lambda);  
}  
generated quantities {  
  real x_ppc;  
  x_ppc <- poisson_rng(lambda);  
}
```

---

## STEP 1: LOOK AT DATA; BUILD POISSON MODEL

```
data {  
  int N;  
  real x[N];  
}  
parameters {  
  real<lower = 0> lambda;  
}  
model {  
  increment_log_prob(cauchy_log(lambda, 0, 1));  
  for (n in 1:N)  
    increment_log_prob(poisson_log(x[n], lambda));  
}  
generated quantities {  
  real x_ppc;  
  x_ppc <- poisson_rng(lambda);  
}
```

---

## STEP 1: LOOK AT DATA; BUILD POISSON MODEL

- ▶ Does the fit look right?
- ▶ Compare the  $x_{ppc}$  to  $x$
- ▶ What needs to be done?

---

## STEP 2: ZERO-INFLATION

- ▶ Add zero inflation.

If  $x[n] == 0$ , then it either came from the Poisson likelihood or it was 0 with probability  $\theta$ .

- ▶ This is a mixture model; also example of marginalizing discrete parameter

$$p(x_n = 0 \mid \lambda, \theta) = \theta \times 1 + (1 - \theta) \times \text{poisson}(x_n \mid \lambda)$$

$$p(x_n \neq 0 \mid \lambda, \theta) = ?$$

- ▶ Make sure to do the posterior predictive check
- ▶ Hint: use the `log_mix` function.



---

## STEP 2: ZERO-INFLATION

- ▶ Add zero inflation.

If  $x[n] == 0$ , then it either came from the Poisson likelihood or it was 0 with probability  $\theta$ .

- ▶ This is a mixture model; also example of marginalizing discrete parameter

$$p(x_n = 0 \mid \lambda, \theta) = \theta \times 1 + (1 - \theta) \times \text{poisson}(x_n \mid \lambda)$$

$$p(x_n \neq 0 \mid \lambda, \theta) = \theta \times 0 + (1 - \theta) \times \text{poisson}(x_n, \lambda)$$

- ▶ Make sure to do the posterior predictive check
- ▶ Hint: use the `log_mix` function.

## STEP 2: ZERO-INFLATION

```
...
parameters {
  real<lower = 0> lambda;
  real<lower = 0, upper = 1> theta;
}
model {
  increment_log_prob(cauchy_log(lambda, 0, 1));
  for (n in 1:N)
    if (x[n] == 0)
      increment_log_prob(log_mix(theta, 0, poisson_log(x[n], lambda);
    else
      increment_log_prob(log1m(theta)
                          + poisson_log(x[n], lambda));
}
generated quantities {
  real x_ppc;
  if (bernoulli_rng(theta) == 1)
    x_ppc <- 0;
  else
    x_ppc <- poisson_rng(lambda);
}
```

---

## STEP 2: ZERO-INFLATION

- ▶ Still wrong. What's going on?

---

## STEP 3: TRUNCATION

- ▶ Add truncation to the model.  
For this exercise, let's fix the upper bound to 9.
- ▶ Now, the likelihood looks like:

$$p(x_n = 0 \mid \lambda, \theta) = \theta \times 1 + (1 - \theta) \times \frac{\text{poisson}(x_n \mid \lambda)}{\text{poisson\_cdf}(U \mid \lambda)}$$

$$p(x_n \neq 0 \mid \lambda, \theta) = (1 - \theta) \times \frac{\text{poisson}(x_n \mid \lambda)}{\text{poisson\_cdf}(U \mid \lambda)}$$

- ▶ Hint: use `poisson_cdf_log` is evaluated on the log scale

## STEP 3: TRUNCATION

```
...
parameters {
  real<lower = 0> lambda;
  real<lower = 0, upper = 1> theta;
}
model {
  increment_log_prob(cauchy_log(lambda, 0, 1));
  for (n in 1:N) {
    real log_trunc_poisson;
    log_trunc_poisson <- poisson_log(x[n], lambda) - poisson_cdf_log(9, lambda);
    if (x[n] == 0)
      increment_log_prob(log_mix(theta, 0, log_trunc_poisson));
    else
      increment_log_prob(log1m(theta) + log_trunc_poisson);
  }
}
generated quantities {
  real x_ppc;
  if (bernoulli_rng(theta) == 1)
    x_ppc <- 0;
  else {
    x_ppc <- poisson_rng(lambda);
    while (x_ppc > 9)
      x_ppc <- poisson_rng(lambda);
  }
}
```

---

## RECAP

- ▶ Wrote model without analytic solution
- ▶ Zero-inflation is example of marginalizing out a discrete parameter
- ▶ Iterated over 3 different models
- ▶ `increment_log_prob()` can be used to write any log joint distribution function

# NON-CENTERED REPARAMETERIZATION


# FUNNEL

►  $y \in \mathbb{R}$

►  $x \in \mathbb{R}^9$

$$p(y, x) = \text{Normal}(y|0, 3)$$

$$\times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$





# WHEN DO YOU SEE THIS?

---

- ▶ Hierarchical models
- ▶ Variance parameters go to 0, all parameters shrink  
Variance parameters get large, all parameters spread
- ▶ Trick to handle low data situations
- ▶ Called non-centered parameterization aka the Matt trick ...

# STEPS

---

1. Add new parameter,  $*_{\text{raw}}$ .
2. Move original parameter to transformed parameters block.
3. Assign transformation of  $*_{\text{raw}}$  to original parameter.
4. Put  $\text{Normal}(0, 1)$  prior on  $*_{\text{raw}}$ .

# CENTERED FUNNEL

---

- ▶ Easy to write in Stan
- ▶ Run. See any problems?

```
parameters {  
    real y;  
    vector[9] x;  
}  
  
model {  
    y ~ normal(0, 3);  
    x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 1

---

- ▶ Add new parameter, `*_raw`.

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 1

---

- ▶ Add new parameter, `*_raw`.

```
parameters {  
  real y;  
  vector[9] x;  
  real y_raw;  
}  
  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 2

---

- ▶ Move original parameter to transformed parameters block.

```
parameters {  
  real y;  
  vector[9] x;  
  real y_raw;  
}  
  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 2

---

- ▶ Move original parameter to transformed parameters block.

```
parameters {  
    vector[9] x;  
    real y_raw;  
}  
transformed parameters {  
    real y;  
}  
model {  
    y ~ normal(0, 3);  
    x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 3

---

- ▶ Assign transformation of  $*_{\text{raw}}$  to original parameter.

```
parameters {  
  vector[9] x;  
  real y_raw;  
}  
transformed parameters {  
  real y;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



# NON-CENTERED FUNNEL: STEP 3

---

- ▶ Assign transformation of  $*_{\text{raw}}$  to original parameter.

```
parameters {  
  vector[9] x;  
  real y_raw;  
}  
transformed parameters {  
  real y;  
  y <- 3 * y_raw;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 4

---

- ▶ Put  $\text{Normal}(0, 1)$  prior on  $*_{\text{raw}}$ .

```
parameters {  
  vector[9] x;  
  real y_raw;  
}  
transformed parameters {  
  real y;  
  y <- 3 * y_raw;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL: STEP 4

---

- ▶ Put  $\text{Normal}(0, 1)$  prior on  $*_{\text{raw}}$ .

```
parameters {  
  vector[9] x;  
  real y_raw;  
}  
transformed parameters {  
  real y;  
  y <- 3 * y_raw;  
}  
model {  
  y_raw ~ normal(0, 1);  
  x ~ normal(0, exp(y/2));  
}
```

# NON-CENTERED FUNNEL

---

- ▶ Repeat for  $x_s$ .
- ▶ Steps:
  1. Add new parameter,  $*_{\text{raw}}$ .
  2. Move original parameter to transformed parameters block.
  3. Assign transformation of  $*_{\text{raw}}$  to original parameter.
  4. Put  $\text{Normal}(0, 1)$  prior on  $*_{\text{raw}}$ .

# NON-CENTERED FUNNEL

---

```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
transformed parameters {  
  real y;  
  vector[9] x;  
  
  y <- 3.0 * y_raw;  
  x <- exp(y/2) * x_raw;  
}  
model {  
  y_raw ~ normal(0, 1);  
  x_raw ~ normal(0, 1);  
}
```

# CENTERED VS NON-CENTERED

---

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
transformed parameters {  
  real y;  
  vector[9] x;  
  
  y <- 3.0 * y_raw;  
  x <- exp(y/2) * x_raw;  
}  
model {  
  y_raw ~ normal(0, 1);  
  x_raw ~ normal(0, 1);  
}
```